

Student critique of AI interpretations

During this assignment, AI was useful for expanding the three selected APIs into systematic domain, security, state-transition, and schema partitions, but its output was not reliable enough to accept unchanged. It produced a syntactically broken JSON idea for LOGIN-035, so I replaced that candidate with an executable empty-object test. For CHECKOUT-034, it initially treated a client-supplied `user_id` as authoritative even though identity must come from the bearer token; for CHECKOUT-035, it assumed idempotency although the specification defines no idempotency key. These errors came from filling gaps with common API conventions instead of distinguishing documented requirements from assumptions. The AI was also incomplete when it treated checkout as an isolated request: it did not initially test an empty cart or verify that the server, rather than the client, owns the order total. Likewise, state-transition generation did not automatically enforce authorization before transition validity, which hid the normal-user role-escalation risk in ORDER-040. Execution was therefore essential. Newman recorded 267 assertions with 50 failures, but those failures were repeated symptoms, not 50 distinct bugs; reviewing case IDs, expected results, and actual responses consolidated them into seven root causes. The most important principle I learned is to treat AI output as a set of test hypotheses, not as evidence. Effective collaboration requires constrained prompts, explicit traceability to the specification and implementation, human review of every expected result, and real execution in a controlled SUT. When the evidence contradicts an AI suggestion, I must preserve the failure, investigate the underlying rule, and correct the model's assumption rather than weakening the assertion.