

HW06 – API Testing Report

Student ID: 23127065

SUT: EShop backend API

SUT commit: d97f995247a4a31ac91e8c6664da6fbf58b5fbd5

Execution date: 2026-08-08

Public repository: <https://github.com/yuran1811/hcmus-sw-testing--hw>

Exact published submission URL: TODO(HUMAN): paste URL after pushing this folder

1. Scope and endpoint selection

The selected endpoints satisfy all three pools and exercise different risk profiles.

Pool	Feature	Endpoint	Primary risks
A	FR-02 login	POST /api/login	credential partitions, enumeration, lockout, sensitive response data
B	FR-08 checkout	POST /api/checkout	authentication, total/address integrity, empty cart, schema
C	FR-18 admin order state	PUT /api/admin/orders/:id/status	role escalation, IDOR, legal/illegal transitions, terminal states

TODO(HUMAN): confirm that no group member selected this same three-API combination.

The API specification was compared with the real Express routes in backend/server.js; the implementation was used for executable setup while the specification and course rules remained authoritative for expected results. This matters because the real backend intentionally contains defects, including a role check absent from the admin route and a canceled-to-delivered transition present only in code.

2. Step-by-step AI-assisted method

The work was intentionally decomposed instead of asking AI for “all API tests” in one prompt.

1. Extract the eight-page assignment and turn each mandatory artifact into an audit checklist.
2. Inspect the published API specification and real routes; select one API from each pool.
3. Query current Postman/Newman documentation with Context7 for collection pre-request headers, environment/data variables, JSON schema checks, reporters, and CI exit behavior.
4. Build domain partitions for every request field, then add authentication/authorization, injection, schema, and state-transition dimensions.
5. Generate 35 candidate cases per endpoint and label each VALID, INVALID, or INCOMPLETE with reasoning and a corrected final expectation.
6. Draft five additional high-risk cases per endpoint. These remain marked Student-extension draft (HUMAN REQUIRED) until the student personally reviews/rewrites them.
7. Export the reviewed matrix to Excel and external JSON; generate a Postman v2.1 collection with a collection-level anti-cheat header script.
8. Run three isolated data-driven Newman folders against a freshly seeded copy of the exact SUT revision. Preserve, classify, and report failures rather than changing assertions to pass.

The complete prompt/output record and declaration are in Appendix A and AI_Audit_Report.md.

3. Test design and human audit

The workbook `test-cases/23127065_HW06_Test_Cases.xlsx` is the source of truth. It has one sheet per API plus a summary sheet. Each test row records ID, origin, initial AI audit label, audit reason, technique, requirement/security/state trace, preconditions, request body, authentication mode, expected status, extra schema assertion, and final expectation.

API	AI candidates	Initial VALID	Initial INVALID	Initial INCOMPLETE	Extension drafts	Final executable rows
Login	35	34	1	0	5	40
Checkout	35	33	1	1	5	40
Order status	35	35	0	0	5	40
Total	105	102	2	1	15	120

Examples of correction include replacing syntactically broken JSON with a transport-valid empty object (LOGIN-035), rejecting an AI assumption that checkout should trust a body `user_id` (CHECKOUT-034), and removing an unspecified idempotency expectation (CHECKOUT-035). The matrix retains the original label/reason so the audit is visible rather than overwritten.

TODO(HUMAN): review all 120 rows and complete the `human_review` column. Rows 036-040 on each sheet must become the student's own reviewed extensions before submission.

4. API 1 – Login pipeline

Generated coverage

The 35 AI candidates cover valid login; whitespace/case behavior; missing, null, empty, numeric, malformed, and oversized email/password values; SQL/XSS payloads; unknown accounts; extra fields; account lockout transitions; JSON content type; generic authentication errors; success schema; and sensitive response fields.

The five extension drafts add whitespace-only credentials, duplicate JSON keys, a prototype-pollution-shaped field, a large irrelevant field, and an explicit anti-cheat-header assertion. AI tends to miss these because ordinary prompts focus on business values rather than parser ambiguity, object-merging hazards, or evidence rules that are external to the API contract.

Execution and findings

Newman ran 40 primary iterations, 91 total requests including setup, and 90 assertions. Seventy assertions passed and 20 failed. The failures map to three genuine defects:

- BUG-API-001: malformed credential domains receive 401 instead of validation 400.
- BUG-API-002: a correct password is blocked after only two failed attempts.
- BUG-API-003: successful login returns the plaintext password in user.

Raw evidence: `newman/results/login-cli.txt`, `login.json`, `login.xml`, and `newman/reports/login/index.html`.

5. API 2 – Checkout pipeline

Generated coverage

The 35 AI candidates cover positive boundaries, zero/negative/decimal/oversized totals, numeric and nonnumeric strings, null/missing/wrong types, address partitions, injection/XSS, mass-assignment-shaped fields,

missing/malformed/tampered tokens, admin-token behavior, empty/array bodies, and success schema/content type.

The five extension drafts focus on business invariants that field-by-field generation missed: rejecting checkout with no cart, refusing a client-supplied total inconsistent with cart contents, an extreme negative value, ignoring a foreign user_id, and numeric parser overflow. AI misses these when the prompt/spec models checkout as an isolated request instead of a server-owned pricing workflow.

Execution and findings

Newman ran 40 primary iterations, 78 total requests including authentication setup, and 91 assertions. Sixty-five assertions passed and 26 failed. The failures consolidate into:

- BUG-API-004: checkout accepts invalid totals, addresses, and body shapes.
- BUG-API-005: checkout creates an order for an empty cart and trusts an arbitrary client total.

Authentication checks themselves behaved as observed: missing token returned 401, and malformed/tampered JWTs returned 403.

Raw evidence: newman/results/checkout-cli.txt, checkout.json, checkout.xml, and newman/reports/checkout/index.html.

6. API 3 – Admin order-state pipeline

Generated coverage

The matrix enumerates all five states: pending, confirmed, shipping, delivered, and canceled. It covers allowed transitions, forbidden skips/backtracking/repeats, terminal-state behavior, invalid enum/type/missing status, nonexistent order, missing/malformed authentication, success schema, injection, role escalation, and IDOR.

The five extension drafts target a normal user changing an admin-managed order, cross-user mutation, exact success schema, status injection, and authorization precedence on a delivered order. AI often misses authorization precedence because a state-machine prompt evaluates transition validity before asking which principal may invoke the transition.

Execution and findings

Newman ran 40 primary iterations, 197 total requests including login/order/state setup, and 86 assertions. Eighty-two assertions passed and four failed:

- BUG-API-006: canceled → delivered returns 200 although canceled is terminal.
- BUG-API-007: normal-user tokens reach and mutate the admin endpoint; authorization is not enforced.

Raw evidence: newman/results/orderstatus-cli.txt, orderstatus.json, orderstatus.xml, and newman/reports/orderstatus/index.html.

7. Postman and Newman features used

Feature	Concrete use/evidence
Collection/folders	One v2.1 collection with three endpoint folders
Collection pre-request script	Upserts and logs X-Student-Id: 23127065 on every primary request
Request pre-request scripts	Register isolated login users; fetch JWTs; create/advance orders
Collection/environment/local variables	studentId, baseUrl, request-scoped order_id
Data-driven runs	One external JSON file per API, 40 iterations each

Feature	Concrete use/evidence
Automated tests	status, header, content type, response shape, sensitive-field checks
CLI and multiple reporters	CLI transcript, JSON, JUnit XML, htmlextra HTML
CI/CD	GitHub Actions workflow with green and intentional-red modes

Context7 was used to verify the current `pm.request.headers.add/upsert` approach, `pm.environment` variables, JSON schema assertions, `newman run -e`, data files, multi-reporter exports, and CI exit semantics. Relevant current documentation: [Postman pre-request scripts](#), [Postman request API](#), and [Newman](#).

Workspace, monitor, and mock-server screenshots are not claimed because they require a real signed-in Postman GUI session. **TODO(HUMAN)**: use and evidence only the additional Postman features you actually perform.

8. CI/CD

The repository workflow `.github/workflows/hw06-api-tests.yml` checks out the exact SUT commit, seeds and starts the API, and runs the pinned local Newman installation. Manual input `pass` executes one known-good case per API. `intentional-failure` executes one deliberately impossible 418 assertion so the red pipeline is unambiguous and not confused with a product defect. Both upload JUnit evidence.

Both modes were pre-verified locally: the `pass` sample produced three green folders with three assertions each, while the `intentional` sample exited nonzero with exactly one failed assertion.

TODO(HUMAN): push and run both modes, then add the two `commit/run` links and screenshots to `ci/CI_CD_Report.md`.

9. Defect reporting

Fifty failed assertions were consolidated by root cause into seven genuine defects rather than filed as 50 duplicates. `Bug_Report.md` contains severity, case IDs, expected/actual results, impact, and evidence paths.

TODO(HUMAN): create seven public GitHub issues and attach a real screenshot to every issue. Paste each URL into `Bug_Report.md`.

10. Agent Skill design and implementation

The reusable skill is in `agent-skill/generate-api-tests/`. Its executable generator exports exactly 35 AI cases plus five extension drafts per endpoint, Excel review sheets, JSON data, Postman collection, environment, and CI samples. It was validated by running the generator and executing its output.

Pseudocode is in `design/pseudocode.md`. The assignment requires a self-drawn, non-AI diagram, so no generated diagram is presented as student work. **TODO(HUMAN)**: add `design/23127065-api-test-generator.png` and describe the drawing tool/date.

Demonstration video: **TODO(HUMAN)**: unlisted video URL.

11. Limitations and reproducibility

- The SUT database was freshly seeded in an isolated temporary copy; the original neighboring checkout was not mutated.
- The committed environment uses port 3000. Local evidence used port 31065 because port 3000 was already occupied; Newman received an explicit `--env-var` override.

- The external repository is public, but new artifacts, GitHub issues, and Actions runs do not exist publicly until the student pushes/publishes them.
- AI-produced audit judgments and extension drafts are not substitutes for the required student review.
- Failing tests are requirement evidence and must not be weakened to produce a green full-suite report.

Appendix A – AI declaration and audit summary

AI was used openly for rubric extraction, documentation retrieval, code/spec comparison, candidate generation, test implementation, report drafting, and local execution analysis. It was not used to fabricate a self-drawn diagram, personal critique, signature, group coordination, screenshots, GitHub issues, CI links, or video. The complete interaction table and declaration are in `AI_Audit_Report.md`; the student must review and sign it before submission.